

De-Slop Stylometrics: A Master Lexical, Syntactic, and Structural Blacklist for Eliminating Synthetic Tells in Large Language Model Generations

Theoretical Foundations of Synthetic Text Detection

Automated artificial intelligence detection architectures—including platforms such as GPTZero, Originality.ai, Copyleaks, Turnitin, and Pangram—do not evaluate semantic intent or factual veracity¹. Instead, these engines rely on statistical stylometrics, analyzing token probability distributions, entropy, and structural variability across sliding text windows².

The primary statistical metrics utilized in AI writing detection are perplexity (PPL) and burstiness (B)². Perplexity measures the log-likelihood unpredictability of word selection relative to an underlying reference language model². Large language models (LLMs) are optimized through next-token prediction, Supervised Fine-Tuning (SFT), and Reinforcement Learning from Human Feedback (RLHF)¹. These training regimes systematically penalize low-probability tokens, causing LLMs to select statistically expected, high-probability words¹. Consequently, synthetic text displays unnaturally low perplexity, generating prose that feels predictable, generic, and mathematically homogenous².

Burstiness quantifies the document-level variance in sentence length, clause density, and rhythmic cadence². Mathematically, burstiness (B) is defined as the standard deviation of sentence word counts (σ) divided by the mean sentence word count (μ):

$$B = \frac{\sigma}{\mu}$$

Human prose exhibits high burstiness, typically scoring between 0.60 and 1.00+, because human thought naturally oscillates between concise declarations, compound analytical structures, rhetorical fragments, and parenthetical interruptions⁶. Conversely, raw LLM outputs cluster within a low burstiness range of 0.15 to 0.30⁶. This statistical uniformity creates a monotonous, metronomic sentence length (typically 15–25 words per sentence) that serves as a primary signature for detection classifiers⁶.

Corpus analysis demonstrates that synthetic vocabulary tells fluctuate dynamically across

model generations and architecture updates⁷. Research from the Max Planck Institute revealed that usage frequencies for terms such as *delve*, *robust*, and *pivotal* spiked by over 50% across published literature following the widespread adoption of instruction-tuned models⁷. Similarly, empirical studies analyzing 15 million biomedical abstracts documented a 654% jump in the usage of the verb *delve* between 2020 and 2023⁷.

Model Generation Shifts in Overrepresented Lexical Sets

GPT-4 Era (2023–Mid 2024)

High density of decorative abstractions (*delve*, *tapestry*, *testament*, *vibrant*, *intricate*, *interplay*, *crucial*, *pivotal*).

GPT-4o Era (Mid 2024–Mid 2025)

Shift toward active corporate participial openers and gerunds (*align with*, *bolster*, *fostering*, *highlighting*, *showcasing*, *underscoring*).

GPT-5 Era (Mid 2025 Onward)

Concentration of conversational softening, participial transitions (*emphasizing*, *enhancing*), and polite organizational framing.

Model-Specific Anomaly (Grok)

Idiosyncratic bias toward pseudo-scientific jargon (*empirical*, *causal*, *correlate*) paired with persistent overuse of underscore.

To systematically bypass detection algorithms and eliminate synthetic tells—frequently categorized as "AI-slop" or "workslop"—content pipelines must enforce strict blacklists covering vocabulary, sentence syntax, formatting conventions, and pragmatic stance¹.

Stylometric Metric	Natural Human Prose Profile	Raw LLM Output Profile (AI-Slop)	Primary Detection Threshold
Perplexity (PPL)	High ($PPL > \quad$); frequent use of unexpected, context-specific tokens ² .	Low ($PPL < \quad$); persistent selection of top-probability token paths ² .	$PPL < \quad$ flags synthetic origin with high confidence ⁶ .

Burstiness (B)	High ($B = 0.60 -$); dynamic mixing of 3-word and 35-word sentences ⁶ .	Low ($B = 0.15 -$); tight clustering around 15–25 words per sentence ⁶ .	$B <$ serves as a primary signal for GPTZero and ZeroGPT ⁶ .
Copulative Density	High; heavy reliance on basic link verbs (<i>is, are, was, were</i>) ¹⁰ .	Low; active suppression of basic link verbs in favor of elevated substitutes ¹⁰ .	Systematic substitution (<i>serves as, stands as</i>) flags AI origin ¹⁰ .
Paragraph Cadence	Variable; ranges from single-line assertions to extended analytical blocks ⁶ .	Formulaic; uniform 3–4 sentence blocks with topic, support, and conclusion ⁷ .	Repeating structural blocks trigger pattern-recognition filters ² .

Comprehensive Lexical Blacklist: Overrepresented AI Vocabulary

Large language models rely heavily on specific vocabulary subsets because safety fine-tuning and RLHF reward polished, non-confrontational, and broadly applicable prose¹. Banning these high-frequency lexical markers forces models to select precise, domain-specific terminology⁷.

High-Frequency Adjectives and Modifiers

Synthetic writing overuses dramatic modifiers to create an artificial sense of importance⁷. These terms weaken prose by substituting hyperbole for substantive detail¹.

Banned AI Adjective	Primary Risk Factor	Recommended Human Alternative
Crucial	Overused indicator of artificial importance ¹	Important, key, essential, necessary ⁷
Pivotal	Formulaic signifier of narrative significance ⁷	Turning point, central, decisive ⁷
Transformative	Unearned claim of	Major, effective, practical,

	profound impact ⁷	structural ⁷
Groundbreaking	Exaggerated innovation claim ⁷	New, novel, unproven, first-of-its-kind ⁷
Robust	Corporate filler adjective (> usage spike) ¹	Strong, durable, reliable, solid ⁷
Seamless	Unrealistic marketing hyperbole ¹	Direct, smooth, simple, easy ⁷
Invaluable	Exaggerated assessment of utility ⁷	Useful, helpful, practical, worth having ⁷
Vibrant	Overused sensory modifier ¹⁰	Active, dense, varied, lively
Intricate / Intricacies	Vague gesture toward structural complexity ¹⁰	Complex, detailed, specific, nuanced
Bespoke / Tailored	Overused marketing jargon ¹	Custom, specific, built-to-order
Ever-evolving	Canned temporal filler ⁷	Changing, fluid, shifting, unstable
Multifaceted	Avoidance of specific categorization ¹⁴	Varied, complex, multi-part
Adept	Generic competency label ¹⁴	Skilled, capable, practiced
Burgeoning	Overused growth indicator ¹⁴	Growing, expanding, rising
Cognizant	Stilted formality indicator ¹⁴	Aware, mindful, conscious
Paramount	Grandiose prioritization term ¹⁴	Main, chief, top, primary
Profound	Unearned claim of depth ⁷	Deep, clear, heavy, serious

Unparalleled	Unverifiable superlative ¹⁴	Rare, exceptional, unmatched
--------------	--	------------------------------

High-Frequency Verbs and Action Terms

Language models favor abstract metaphorical verbs over concrete actions¹. These verbs sound proactive but often obscure specific mechanisms¹.

Banned AI Verb	Primary Risk Factor	Recommended Human Alternative
Delve / Delving	654% usage spike; primary detection flag ¹	Examine, study, look into, dig into ⁷
Leverage	Excessive utility verb ⁷	Use, deploy, apply, turn to ⁷
Harness	Pseudo-technical utility action ⁷	Use, direct, channel, put to work ⁷
Foster	Over-indexed relational verb ⁷	Build, encourage, support, grow ⁷
Optimize	Generic efficiency claim ⁷	Improve, refine, tune, adjust ⁷
Facilitate	Administrative abstraction ⁷	Help, run, ease, set up ⁷
Elevate	Vague value-proposition term ¹	Raise, boost, improve, lift
Underscore	Overused analytical indicator ⁷	Show, highlight, point to, make clear ⁷
Showcase	Decorative visual metaphor ¹⁰	Present, display, show ¹¹
Augment	Technical padding verb ¹⁴	Add to, expand, supplement
Embark (on a journey)	Cliché narrative opener ¹⁴	Start, begin, set out

Unlock / Unleash	Formulaic marketing metaphor ¹	Access, open, release, enable
Cultivate	Overused development metaphor ⁷	Build, grow, form
Elucidate	Stilted academic verb ¹⁴	Explain, clarify, spell out
Navigate	Metaphorical substitute for management ⁷	Deal with, handle, work through ⁷
Resonate	Overused engagement indicator ¹⁴	Connect, strike home, make sense
Streamline	Corporate optimization filler ¹⁴	Simplify, cut, trim, speed up
Transcend	Grandiose boundary metaphor ¹⁴	Go beyond, exceed, surpass

High-Frequency Nouns, Metaphors, Transitions, and Corporate Workslop

Language models rely on predictable abstract nouns, decorative metaphors, and rigid transitional phrases to connect ideas¹. Eliminating these filler terms forces prose to move directly between concrete points⁷.

Category	Blacklisted AI Words and Exact Phrases	Recommended Human Alternative
Abstract Metaphors	Tapestry, landscape, realm, beacon, cornerstone, treasure trove, linchpin, kaleidoscope, epicenter, cornerstone ¹	System, area, field, sector, core part, collection ⁷
Conceptual Nouns	Testament, paradigm, journey, synergy, plethora, multitude, game-changer, touchpoint, North Star ¹	Proof, sign, method, process, excess, collaboration ⁷

Formal Transitions	Moreover, Furthermore, Additionally, Consequently, Nevertheless, Notably, Ultimately, Thereby ²	Also, plus, and, so, still, but, in the end ⁷
Corporate Workstop	Actionable insights, dynamic environment, strategic alignment, operational excellence, capacity building, thought leadership, value proposition, cost optimization ¹	Useful data, fast setting, clear goals, strong performance, training
Preambles & Openers	In an ever-evolving landscape, In today's fast-paced world, In the realm of, Navigating the complexities of, When it comes to ¹	[Delete entirely; state the thesis immediately] ⁷
Empty Fillers	It is important to note, It is worth noting that, Broadly speaking, At the end of the day, Cannot be overstated ⁷	[Delete completely or replace with direct assertion] ⁷

Syntactic and Structural Pattern Blacklist

Exclusively replacing individual words is insufficient to bypass modern detection models if the underlying sentence mechanics remain formulaic¹². Detection engines evaluate syntactic predictability, structural symmetries, and punctuation distributions¹.

The Copulative Avoidance Trap

A key structural tell of synthetic text is copulative verb avoidance¹⁰. Language models systematically replace simple link verbs (*is, are, was, were*) with inflated functional substitutes to make sentences sound more authoritative¹⁰.

- **Synthetic Substitution:** *"The initiative serves as a testament to organizational resilience and operates as a catalyst for growth."*
[cite: 10, 14]
- **Human Restructure:** *"The initiative shows the organization's resilience and drives growth."*

Blacklisted copulative replacement verbs include¹⁰:

- *serves as / stands as / marks*

- *functions as / operates as / represents [a]*
[cite: 10]
- *boasts / features / maintains / offers [a]*
[cite: 10]
- *underscores / highlights / reflects [a]*
[cite: 9, 10]

Formulaic Binary Parallelisms and Rhetorical Templates

Language models rely on balanced binary contrasts to simulate analytical depth¹. These repetitive structures create distinct mathematical signatures in detection models¹.



- **Example 1:** Banned: *"It's not just about efficiency; it's about transformation."*
[cite: 1]
 - Correction: *"Efficiency matters, but long-term success requires structural change."*
- **Example 2:** Banned: *"No theory. No fluff. Just execution."*
[cite: 1]
 - Correction: *"The plan focuses purely on execution."*
- **Example 3:** Banned: *"And the solution? Better brief alignment."*
[cite: 1]
 - Correction: *"Aligning creative briefs resolves this issue."*

Punctuation Anomalies and Formatting Artifacts

The misuse of specific punctuation marks and layout structures immediately flags synthetic text⁸.

- **Em Dash Overuse (—):** LLMs over-index on em dashes to append explanatory summaries or rhetorical flourishes to sentence endings⁸. Em dash usage should be restricted to a maximum of one occurrence per 1,000 words⁸.
- **The Rule of Three (Triadic Clustering):** LLMs naturally default to rhythmic triads of adjectives, nouns, or clauses (e.g., *"Focused. Aligned. Measurable."* or *"fostering innovation, driving growth, and elevating value"*)¹. Content engineers must break these

triads by reducing items to pairs or expanding them into independent sentences⁸.

- **Boldface Lead-in Vertical Lists:** AI tools rely heavily on bulleted lists where every item begins with a bolded 2–4 word title followed by a colon¹⁰. These vertical structures should be converted into flowing narrative prose⁷.
- **Metronomic Paragraph Structure:** Default LLM generations exhibit a rigid paragraph architecture consisting of one introductory topic sentence, two supporting detail sentences, and a summary sentence⁷. Human writing varies paragraph lengths dynamically, mixing brief single-sentence assertions with extended analytical blocks⁶.

Syntactic Structure	Primary Structural Risk	Humanizing Restructure Directive
Copulative Substitution (<i>"serves as a reminder"</i>)	Eliminates direct subject-verb links; inflates word count ¹⁰ .	Restore direct copulatives (<i>"is a reminder"</i>) or active verbs (<i>"reminds"</i>) ¹⁰ .
Negative Parallelism (<i>"Not X, but Y"</i>)	Highly predictable binary sentence structure ¹⁰ .	State the positive claim directly without the binary contrast ¹¹ .
Rhetorical Question & Answer (<i>"The fix? X."</i>)	Synthetic conversational gimmick ¹ .	Merge into a direct declarative sentence (<i>"The fix requires X."</i>) ¹ .
Triadic Adjective Stacking (<i>"Robust, scalable, agile"</i>)	Predictable rhythm trigger for burstiness detectors ¹ .	Keep a single precise adjective or expand into detailed prose ⁷ .
Bold Lead-in Bullet Lists	Standard structural artifact of instruction tuning ¹⁰ .	Convert bullet points into flowing narrative paragraphs ⁷ .
Em Dash Appends (<i>"text—showing clear value"</i>)	Overused explanatory tail-piece ⁸ .	Split into two distinct sentences or use commas ⁸ .

Pragmatic, Semantic, and Stance Anti-Patterns

Language models exhibit characteristic pragmatic weaknesses stemming from safety alignment and instruction-following protocols¹. These training constraints induce excessive hedging, missing personal perspective, and over-reliance on formulaic conversational wrappers¹.

Epistemic Hedging and Vague Attribution

To maintain neutrality and avoid absolute claims, LLMs insert non-committal qualifiers that erode authoritative tone⁷:

- **Epistemic Hedges:** *It is important to note that, it can be argued that, one might suggest that, it is worth considering that, broadly speaking, generally speaking*⁷.
- **Weasel Wording & Vague Authority:** *Observers have cited, experts argue, industry reports suggest, critics contend, several publications note* (when serving as generic attributions without concrete citations)¹⁰.

Human prose takes clear, defensible stances and attributes claims directly to named, verified sources⁷.

Conversational Preambles and Meta-Commentary

Language models routinely output introductory pleasantries, prompt-acknowledgments, and meta-commentary that delay direct engagement with the subject matter¹.

- **Chatbot Preambles:** *Certainly! Here is a detailed analysis of..., Sure thing! Let's explore..., Based on the information provided...*¹⁴.
- **Introductory Throat-Clearing:** *In today's fast-paced digital world..., In the realm of modern technology..., Welcome to the world of..., When it comes to understanding X...*¹.
- **Self-Reflective Meta-Commentary:** *In this section, we will explore..., As mentioned previously..., This guide will delve into...*¹.

Content pipelines must strip out meta-commentary entirely, starting the draft directly with the core thesis or immediate data point⁷.

Symmetrical Summaries and Formulaic Conclusions

Default LLM generations almost always close with symmetrical, high-level summaries that rehash earlier points without offering new analysis⁷.

Earned Concluding Structural Tropes

Formal Summary Openers: "In conclusion," "In summary," "To summarize," "Ultimately," "At the end of the day," (cite: 1, 7, 8, 14)

High-Level Synthesis Flags: "This serves as a testament to..." "As we look to the future..." "will continue to play a vital role..."

Balanced Wrap-ups: A restatement of the prompt + high-level acknowledgment of challenges + optimistic closing sentence

Human writers close arguments naturally once the core points are established—ending on a concrete observation, a key limitation, or a specific call to action, rather than relying on formulaic summary markers⁷.

System-Level Anti-AI Prompt Engineering Directives

To consistently produce human-like output from LLMs, blacklists and structural guidelines must be integrated directly into system prompts, custom instructions, or model style profiles⁸.

System-level instructions effectively enforce compliance across long generation contexts¹¹.

Master System Prompt Directive Blueprint

The following structured directive block can be integrated into system prompts across platforms (ChatGPT Custom Instructions, Claude System Prompts, Gemini System Instructions, or API configurations) to systematically suppress synthetic tells⁸:

SYSTEM CONFIGURATION: STYLOGRAPHIC HUMANIZATION ENGINE

ROLE: You are an elite domain copy editor and prose architect. Your objective is to write authoritative, natural text completely free of synthetic AI patterns, conversational fluff, and low-perplexity markers.

1. LEXICAL BLACKLIST (STRICT COMPLIANCE)

- **NEVER USE THE FOLLOWING WORDS:** delve, tapestry, landscape, realm, robust, pivotal, crucial, leverage, harness, foster, elevate, underscore, showcase, testament, beacon, bespoke, tailored, intricate, vibrant, transformative, groundbreaking, seamless, invaluable, plethora, multitude, game-changer, touchpoint, burgeoning, paramount.
- **NEVER USE CORPORATE WORKSLOP:** actionable insights, dynamic environment, strategic alignment, operational excellence, capacity building, thought leadership, value proposition, cost optimization, ever-evolving, in today's world.

2. SYNTACTIC CONSTRAINTS & BURSTINESS ENFORCEMENT

- **TARGET BURSTINESS** ($B > 0.70$): You MUST vary sentence length continuously. Alternate short assertions (3–6 words) with compound analytical sentences (25–35 words). Avoid metronomic 15–20 word sentence patterns.
- **ENFORCE DIRECT COPULATIVES:** Use basic link verbs (*is, are, was, were*). NEVER substitute them with *serves as, stands as, functions as, operates as, boasts, or offers*.
- **STRICT SVO SYNTAX:** Write primarily in clear Subject-Verb-Object order.
- **BANNED RHETORICAL STRUCTURES:**
 - Do NOT use "Not only X, but also Y" or "Not X, but Y" constructions.
 - Do NOT use "No X. No Y. Just Z." templates.
 - Do NOT use "It's not about X, it's about Y." templates.
 - Do NOT use rhetorical question-and-answer setups ("And the fix? X").
- **PUNCTUATION RESTRICTION:** Absolute ban on em dashes (—). Limit hyphens. Avoid rule-of-three triads.

3. STANCE & PRAGMATICS

- **ZERO HEDGING:** Eliminate *it is important to note, one might suggest, it can be argued, and broadly speaking*. Make direct, defensible assertions.
- **ZERO THROAT-CLEARING:** Delete greetings, introductions, and meta-commentary ("Certainly!", "In this article..."). Start directly with the core thesis.

- **NO FORMULAIC CONCLUSIONS:** Do NOT use *in conclusion*, *in summary*, or optimistic wrap-up boilerplate. Stop writing immediately after the final argument is delivered.

Comparative Domain Transformation Matrix

Applying systematic de-slopping rules transforms synthetic AI output into concise, authoritative human prose across diverse writing contexts¹.

Domain / Context	Raw Synthetic Generation (AI-Slop Output)	Humanized De-Slopped Transformation
B2B SaaS Marketing	<i>"In today's fast-paced digital landscape, leveraging cloud-based solutions is crucial to unlock robust efficiency and foster seamless collaboration. Our state-of-the-art platform serves as a testament to innovation, empowering decision-makers to elevate operational excellence."</i> [cite: 1, 7, 14, 15]	<i>"Cloud software cuts team coordination costs. We built our platform to automate routine database tasks, letting engineering managers ship code faster without manual infrastructure tracking."</i>
Academic & Historical	<i>"The subject's pivotal role in the revolution stands as a crucial milestone. Furthermore, delving into the intricacies of his legacy underscores a vibrant tapestry of leadership, showcasing his enduring impact."</i> [cite: 7, 9, 10, 14]	<i>"He led three tactical retreats during the 1789 uprising. His command sustained army ranks through winter supply shortages, keeping the core fighting force intact until foreign aid arrived."</i>
Corporate Strategic Brief	<i>"Navigating the complexities of market penetration requires a holistic approach. It is important to note that success is not just about aggressive sales, but also strategic alignment. In</i>	<i>"Entering the European market requires local compliance clearing before hiring sales teams. Regulatory delays average six months; signing distributor contracts early offsets those initial</i>

	<i>conclusion, adopting best practices will drive impact."</i> [cite: 1, 7, 10, 14]	<i>overhead costs."</i>
Technology Product Review	<i>"When it comes to hardware, this device functions as a game-changer. Featuring a robust chassis and seamless software integration, it promises to elevate user experience to new heights while fostering productivity."</i> [cite: 1, 7, 10, 14]	<i>"The laptop's aluminum frame resists flexing under heavy pressure. The updated display drivers fix previous sleep-wake latency, adding roughly forty minutes of active battery life."</i>

Strategic Implementation Roadmap and Conclusion

Eliminating synthetic writeslop requires a structured, multi-pass workflow combining system-level prompt constraints, post-generation editing, and ongoing style profiling⁸. Replacing vocabulary alone changes short-term probability signals but often leaves underlying sentence structures intact, allowing updated detection models to flag the text¹². Complete de-slopping relies on addressing lexical, syntactic, and structural patterns simultaneously⁶. The implementation roadmap follows three core operational phases:

First, content engineering teams must integrate system-level directives containing strict lexical and syntactic blacklists directly into target LLMs¹¹. Establishing custom voice profiles via systematic instructions grounds outputs in human cadence before text generation begins⁸. Second, raw outputs must go through an automated or manual editing pass designed to identify residual AI tells¹¹. Editors should strip meta-commentary, eliminate copulative substitutions (*serves as*, *stands as*), break triadic adjective stacks, remove em dashes, and delete generic conclusion wrappers⁷.

Third, content engineers must actively manage sentence length variation to maintain human-like burstiness ($B > 0.70$)⁶. Enforcing direct Subject-Verb-Object (SVO) structures, mixing brief declarative statements with compound analytical clauses, taking clear stances, and removing epistemic hedges ensures generated text reads naturally and withstands statistical AI detection algorithms².

Works cited

1. List of 300+ AI Words, Phrases and Sentences to Avoid (2026) - Content Beta, <https://www.contentbeta.com/blog/list-of-words-overused-by-ai/>

2. How AI Detectors Work: Perplexity & Burstiness 2026 - AI Free Text Pro,
<https://aifreetextpro.com/blog/how-ai-detectors-work>
3. Why Perplexity and Burstiness Fail to Detect AI | Pangram Labs,
<https://www.pangram.com/blog/why-perplexity-and-burstiness-fail-to-detect-ai>
4. AI Plagiarism Detection — GPTZero Reverse Engineering - Medium,
<https://medium.com/@tigerryan21/ai-plagiarism-detection-gptzero-reverse-engineering-85724f3f0695>
5. Detecting the corruption of online questionnaires by artificial intelligence - PMC,
<https://pmc.ncbi.nlm.nih.gov/articles/PMC10869497/>
6. What Is Burstiness in AI Detection? Complete Technical Explanation [2026],
<https://thehumanizeai.pro/articles/what-is-burstiness-ai-detection-explained>
7. 50 Words AI Overuses (And What to Write Instead) - HumanizeThisAI,
<https://humanizethisai.com/blog/50-words-ai-overuses>
8. Ruben Hassid (@ruben): "How to avoid sounding like an AI, forever: (worth saving for later, I promise) More here: <http://ruben.substack.com>. Taken from: <https://ruben.substack.com/p/detection> Archive of free guides: <https://docs.google.com>, <https://substack.com/@ruben/note/c-186925761>
9. Wikipedia:WikiProject AI Cleanup/Guide,
https://en.wikipedia.org/wiki/Wikipedia:WikiProject_AI_Cleanup/Guide
10. Wikipedia:Signs of AI writing,
https://en.wikipedia.org/wiki/Wikipedia:Signs_of_AI_writing
11. humanize-text · GitHub Topics, <https://github.com/topics/humanize-text>
12. What's the best humanizer skill out there? : r/claudekills - Reddit,
https://www.reddit.com/r/claudekills/comments/1v8wa5r/whats_the_best_humanizer_skill_out_there/
13. THE RUBEN HASSID BLUEPRINT: Complete Creator & Channel Analysis - buldrr ai,
<https://buldrr.com/the-ruben-hassid-blueprint/>
14. <https://embryo.com/blog/list-words-ai-overuses/>
15. Gemini Prompt Library - Cuppa,
<https://learn.cuppa.ai/prompting-guidelines-and-library/gemini-prompt-library>
16. How Students Bypass AI Detection: Risks and Reality - Lynote.ai,
<https://lynote.ai/blog/how-students-bypass-ai-detection>
17. How To Identify AI Text In 2026: The Updated Map - Headcore Digital,
<https://headcore.digital/en/how-to-identify-ai-text-2026-updated-map-patterns>
18. anti-ai-writing · GitHub Topics, <https://github.com/topics/anti-ai-writing>